

[CLASSIFIED · PRELIMINARY ROUND]

THE CORPORATE HEIST.

*Two hours. Ten thousand profiles. Five hundred names that matter.
A problem statement for the online preliminary round.*

FIELD	DETAIL
DATE	Saturday, 19 September 2026
WINDOW	6:00 PM to 8:00 PM IST. The deadline is strict.
CHANNEL	Unstop: problem statement, dataset link, submissions, all communication
TEAMS	2 to 5 members, online
DELIVERABLE	One zip on Unstop: ranked shortlist, code, documentation
STAKES	Top 40 teams advance to the finals at IIT Delhi, 28 September 2026

READ THIS FILE IN FULL BEFORE YOU WRITE ANY CODE. QUERIES ON UNSTOP ONLY. THERE ARE NO PRIVATE HINTS.

FILE 01 // BACKGROUND

01 THE BRIEF

Nightingale Systems is one of the most secretive technology employers in the country. It never publishes how it hires, and its offer letters go out without warning. Your team works for **Kestrel Talent**, a rival agency whose business depends on reaching the right people before Nightingale does.

Last night, an insider walked out of Nightingale's offices with three things.

EXHIBIT	CONTENTS
A • THE ARCHIVE	Nightingale's historical hiring records: 20,000 people hired in past cycles, each with the performance score they earned after one year in the job.
B • THE LEDGER	A sealed record from one past cycle: that cycle's applicant pool, and the IDs of the applicants who went on to become Nightingale's top 5%.
C • THE VAULT	This cycle's applicant pool: 10,000 profiles, with no outcomes. Nightingale makes its offers on Monday.

Your mission is to **identify the 500 applicants in the Vault who will become Nightingale's top 5% this cycle**, and rank them, so that Kestrel can reach them first.

1.1 THE INSIDER'S DEBRIEF

Before disappearing, the insider left a voice note. Nightingale's own records do not explain themselves; this note is the only first-hand account you have of how things work inside.

TRANSCRIPT • VOICE NOTE, RECOVERED		UNEDITED • 1 SPEAKER
00:03	Don't assume this cycle works like the old ones. Things changed after the reorg.	
00:11	They finally started logging public code contributions this year. You'll only see that field in the new extract; the old records never had it. The managers are obsessed with it. A handful of contributions means nothing to them, but someone with a sustained record, dozens of merged contributions or more in the year, gets fast-tracked, even if the rest of the profile is only good rather than great.	
00:42	The hiring committee got burned last year. Résumés where the titles climb faster than the years behind them, 'Head of' something three years out of college, now go straight to the bin, however shiny the rest looks.	
01:05	Headcount has to be in seats before the quarter closes. Anyone who can't join inside two months is dead to them this cycle. Doesn't matter how good they are.	
01:21	And the thing nobody outside knows: their best people this cycle came from colleges they had never hired from before. Candidates who topped the technical assessment and genuinely fit the role. No pedigree, no problem. They love those.	

01:39

And forget the old committee's pet preferences. Big-name colleges, candidates from the big metros, glowing referrals, big-brand employers, 'proper' degrees, CVs without a single gap: the old panel quietly inflated all of it. The new panel ignores every one of those. The only old habit that survived the reorg is the old-boys' network: a handful of colleges where Nightingale's own leadership studied still get a leg up, same as always.

02:07

One more thing. The data is filthy. Recruiters type whatever they want, however they want. Some profiles are straight-up fabricated: they look amazing and don't add up. And the same person sometimes shows up twice because two different recruiters entered them. Nightingale would never hire a fake, and it only hires a person once.

What this means for you. The Archive and the Ledger describe how Nightingale hired in the past. The Vault is this cycle. Your shortlist must reflect what Nightingale will actually do this cycle.

FILE 02 // OBJECTIVE

02 THE TASK

Using the Archive, the Ledger and the debrief, produce a ranked shortlist of **exactly 500 candidate_id values from test.csv** (the Vault): the applicants you believe Nightingale will rate as its top 5% this cycle, ordered from most to least likely.

- **Rank 1** is the applicant you are most confident about.
- A fabricated profile on your shortlist, or the same person listed twice under different IDs, is a failure of the shortlist, not a harmless extra.
- You are not predicting a score for every applicant. Only the membership and the order of your 500 matter.

FILE 03 // EVIDENCE

03 THE DATA

A single zip is released on Unstop at **6:00 PM IST**. It contains the following.

FILE	ROWS	WHAT IT IS
train.csv	20,000	The Archive. Historical hires with the outcome column <code>post_hire_score</code> (0-100), the one-year performance score. Scores are noisy, as real performance reviews are.
dev.csv	2,999	The Ledger pool. Same columns as <code>train.csv</code> , without <code>post_hire_score</code> . Drawn from a past cycle.
dev_winners.csv	150	The Ledger's outcome: rank and <code>candidate_id</code> of the applicants in <code>dev.csv</code> who became that cycle's top 5%. Each real person appears at most once. Use it to check your approach end to end.
test.csv	10,000	The Vault: this cycle's pool. No outcomes. Contains one column the older records do not have.

FILE	ROWS	WHAT IT IS
sample_submission.csv	500	The exact submission format. Its contents are arbitrary.
check_format.py	–	Verifies that your submission.csv is well formed. It does not score anything.
documentation_template.md	–	The structure your documentation must follow.
Corporate_Heist_...pdf	–	This document.

3.1 COLUMN REFERENCE

All candidate files share the columns below; `post_hire_score` is only in `train.csv` and `public_code_contributions` only in `test.csv`. Values are shown as recruiters typed them: formats, units and spellings vary within a column. That is deliberate.

COLUMN	DESCRIPTION
candidate_id	Unique ID of a row (CH-XXXXXX). A row is a profile entry, not necessarily a unique person.
full_name, email, phone	Contact details as entered by the recruiter.
age	Age in years at the time of application (occasionally blank).
graduation_year	Year the highest degree was completed.
institute	Institute of the highest degree, as typed by the recruiter.
degree, major	Degree name and discipline, as typed.
current_city	City the candidate currently lives in.
applied_role	The role applied for at Nightingale (10 roles).
current_title	Current job title.
career_path	Chronological title history, with time spent in each title.
total_experience	Total professional experience.
num_employers	Number of employers worked for.
company_type, company_size	Type and size of the current employer.
skills	Skills listed on the profile.
certifications	Professional certifications.
technical_assessment	Score on Nightingale's role-specific technical assessment (out of 100). Blank or marked when not taken.
aptitude_score	Score on the general aptitude test (out of 10).
last_rating	Rating in the most recent performance cycle at the current employer (1-5).
kpi_met	Whether the candidate met more than 80% of their KPIs last year.
awards	Awards won in the last year.

COLUMN	DESCRIPTION
trainings_last_year, training_hours	Trainings attended last year; total training hours logged.
recruitment_channel	How the candidate entered the pipeline.
last_job_change	Years since the last job change.
currently_enrolled	Whether the candidate is currently enrolled in a course.
overtime_history	Whether the candidate routinely worked overtime in the current job.
current_ctc, expected_ctc	Current and expected annual compensation.
notice_period	How soon the candidate can join.
recruiter_note	Free-text note written by the recruiter after the screening conversation.
public_code_contributions	test.csv only. Public code contributions in the last 12 months, as recorded by Nightingale's recruiters this cycle.
post_hire_score	train.csv only. The hire's one-year performance score (0-100). Your training target.

3.2 FACTS YOU CAN RELY ON

- The data is synthetic. It was built from several public HR datasets that were mixed, re-sampled and substantially transformed; the outcomes were generated independently of those sources. Looking up the original datasets will not reveal any outcome.
- Protected attributes such as gender have been removed from every profile.
- A person can appear more than once, under different candidate_id values, with cosmetic differences in how their details were entered.
- Every row in all three files is a real applicant record in the story. None of the rows are padding.
- The Archive and the Ledger were produced under the same, older hiring standard. The Vault is this cycle.

FILE 04 // DELIVERY

04 WHAT TO SUBMIT

Upload **one zip file of at most 20 MB** to the Unstop round before **8:00 PM IST**. You may re-upload as often as you like; only your last upload before the deadline counts.

IMPORTANT

All three parts are required. The round description on Unstop mentions documentation; for this round the zip must contain your shortlist (submission.csv), your code and your documentation. A zip without a valid submission.csv cannot be assessed.

Upload early. Aim to have a valid zip uploaded by 7:30 PM and overwrite it if you improve. Uploads still in progress at 8:00 PM are not accepted, and no submissions are taken by email or chat.

The zip must have exactly this layout:

```

<your_team_name>.zip
├── submission.csv
├── documentation.pdf
└── code/
    ├── main.py          (or main.ipynb)
    └── ...              helper files, small external data
  
```

4.1 SUBMISSION.CSV

- Header row exactly rank,candidate_id.
- Exactly 500 data rows. rank runs 1 to 500 with no gaps; each candidate_id appears once and exists in test.csv.
- Run `python check_format.py submission.csv test.csv` before uploading. A shortlist that is not well formed may not be assessable.

4.2 CODE/

- Your code must reproduce your submission.csv. It is run in a clean environment with **no internet access**, with train.csv, dev.csv, dev_winners.csv and test.csv in the working directory, and must write submission.csv there.
- Entry point: code/main.py or a single code/main.ipynb, run top to bottom.
- Limits: CPU only, 2 cores, 4 GB RAM, **5 minutes** wall-clock. Set all random seeds.
- Use only the libraries in Appendix A. Any external data your code reads must be inside the zip.
- Do not include train.csv, dev.csv, test.csv or large model files.

4.3 DOCUMENTATION.PDF

- At most 4 pages, following documentation_template.md (see Appendix B).
- It must describe what your code actually does. It is read alongside your code.

PRE-UPLOAD CHECKLIST

- [] `check_format.py` prints OK for my `submission.csv`.
- [] My code, run from a clean folder, regenerates the same `submission.csv` within 5 minutes.
- [] `documentation.pdf` is 4 pages or fewer and declares every external resource and AI tool used.
- [] The zip is under 20 MB and follows the layout above.

FILE 05 // RULES

05 RULES OF ENGAGEMENT

5.1 ALLOWED

- **AI assistants** (ChatGPT, Claude, Copilot and similar) for writing code and analysis. Declare them in your documentation.
- **Public external data, libraries and pretrained models** (for example institute lists, skill taxonomies, currency references). Each must be publicly available, cited in your documentation, and, if your code reads it, included in your zip.
- Any modelling approach, from hand-written rules to machine learning.

5.2 NOT ALLOWED

- **Sharing or receiving solutions, shortlists or code between teams**, in any form, during or after the window. Both the team that shares and the team that receives are disqualified.
- Submitting on behalf of another team, or one person competing in more than one team.
- Private or paid data obtained specifically for this competition.
- Any attempt to access, probe or interfere with the organisers' systems.

5.3 INTEGRITY

- Every submission is checked for integrity. Code that does not reproduce its shortlist, documentation that does not match the code, or shortlists that match other teams' shortlists beyond what independent work can explain will be reviewed and may lead to disqualification. Flagged teams are contacted before any decision is made.
- **Day-1 audit at IIT Delhi.** On 28 September, every qualifying team explains its preliminary-round code in person to Eightfold.ai engineers before the finals begin. A team that cannot explain its own submission is disqualified and replaced from the waitlist.
- The organisers' decisions on scoring, integrity and qualification are final.

FILE 06 // VERDICT

06 HOW SHORTLISTS ARE JUDGED

Each shortlist is compared with what actually happened: the applicants in the Vault who became Nightingale's top 5% this cycle. Qualification is decided by the quality of your shortlist, subject to

the integrity checks in 5.3. The top 40 teams qualify for the finals at IIT Delhi on 28 September 2026 (RDV Day 1); a waitlist is maintained.

FILE 07 // SCHEDULE

07 TIMELINE

WHEN (IST)	WHAT HAPPENS
19 SEP • 6:00 PM	Round opens on Unstop. This document and the dataset link are published.
6:00 PM – 8:00 PM	Build and upload. Organisers answer queries on Unstop throughout.
19 SEP • 8:00 PM	Hard deadline. The Unstop upload closes. No late submissions for any reason.
AFTER EVALUATION	Qualified teams (top 40) and the waitlist are announced on Unstop.
WITHIN 48 H	Qualified teams confirm attendance, or the place passes to the waitlist.
28 SEP • 10 AM – 4 PM	Finals at IIT Delhi, RDV Day 1. The Day-1 audit comes first.

FILE 08 // COMMUNICATION

08 QUERIES

All communication for this round happens on Unstop. Raise any query on the event's Unstop page and the organisers will respond there. Answers of general interest are shared with all teams, so everyone works from the same information.

Organisers will clarify the rules, the files and the submission process. They will not comment on approaches, data patterns, or whether a particular idea is right.

FILE A // APPENDIX

A LIBRARIES

When your code is re-run, it has Python 3.11 with the following, and nothing else beyond the standard library. Nothing can be downloaded.

```
numpy · pandas · scipy · scikit-learn · xgboost · lightgbm · catboost  
statsmodels · rapidfuzz · sentence-transformers  
  
pre-loaded model:  all-MiniLM-L6-v2    (load it by this name)
```

FILE B // APPENDIX

B DOCUMENTATION TEMPLATE

documentation_template.md in the zip has these sections. Keep to them, and to 4 pages.

- 01 Summary
- 02 Data cleaning and parsing
- 03 What drives a great hire: your findings
- 04 The current cycle: what differs in the Vault, and how your shortlist accounts for it
- 05 Profiles you excluded, and why
- 06 Model and selection procedure
- 07 What did not work
- 08 External resources and AI tools (mandatory)
- 09 How to run

— END OF FILE —